

```
In [1]: from pydoc import help
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import scale
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from scipy import stats
from IPython.display import display, HTML
%matplotlib inline
np.set_printoptions(suppress=True)
DISPLAY_MAX_ROWS = 20
pd.set_option('display.max_rows', DISPLAY_MAX_ROWS)
```

```
In [3]: data = pd.read_csv("http://archive.ics.uci.edu/ml/machine-learning-databases/w
data.columns = ["V"+str(i) for i in range(1, len(data.columns)+1)]
data.V1 = data.V1.astype(str)
X = data.loc[:, "V2":]
y = data.V1
data
```

```
Out[3]:
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	5.64	1.04	3.92	1065
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	3.40	1050
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	3.17	1185
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	7.80	0.86	3.45	1480
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	4.32	1.04	2.93	735
...
173	3	13.71	5.65	2.45	20.5	95	1.68	0.61	0.52	1.06	7.70	0.64	1.74	740
174	3	13.40	3.91	2.48	23.0	102	1.80	0.75	0.43	1.41	7.30	0.70	1.56	750
175	3	13.27	4.28	2.26	20.0	120	1.59	0.69	0.43	1.35	10.20	0.59	1.56	835
176	3	13.17	2.59	2.37	20.0	120	1.65	0.68	0.53	1.46	9.30	0.60	1.62	840
177	3	14.13	4.10	2.74	24.5	96	2.05	0.76	0.56	1.35	9.20	0.61	1.60	560

178 rows × 14 columns


```
In [4]: data.loc[:, "V2":"V6"]
```

```
Out[4]:
```


	V2	V3	V4	V5	V6
0	14.23	1.71	2.43	15.6	127
1	13.20	1.78	2.14	11.2	100
2	13.16	2.36	2.67	18.6	101
3	14.37	1.95	2.50	16.8	113
4	13.24	2.59	2.87	21.0	118
...
173	13.71	5.65	2.45	20.5	95
174	13.40	3.91	2.48	23.0	102
175	13.27	4.28	2.26	20.0	120
176	13.17	2.59	2.37	20.0	120
177	14.13	4.10	2.74	24.5	96

178 rows × 5 columns

```
In [7]: pd.plotting.scatter_matrix(data.loc[:, "V2":"V6"], diagonal="kde")
plt.tight_layout()
plt.show()
```



```
In [6]: ax = data[["V2", "V3", "V4", "V5", "V6"]].plot()  
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5));
```



```
In [8]: X.apply(np.mean)
```

```
Out[8]: V2      13.000618  
V3       2.336348  
V4       2.366517  
V5      19.494944  
V6      99.741573  
V7       2.295112  
V8       2.029270  
V9       0.361854  
V10      1.590899  
V11      5.058090  
V12      0.957449  
V13      2.611685  
V14     746.893258  
dtype: float64
```

```
In [9]: X.apply(np.std)
```

```
Out[9]: V2      0.809543
        V3      1.114004
        V4      0.273572
        V5      3.330170
        V6     14.242308
        V7      0.624091
        V8      0.996049
        V9      0.124103
        V10     0.570749
        V11     2.311765
        V12     0.227929
        V13     0.707993
        V14    314.021657
        dtype: float64
```

```
In [10]: class2data = data[y=="2"]
        class2data.loc[:, "V2"].apply(np.mean)
        class2data.loc[:, "V2"].apply(np.std)
```

```
Out[10]: V2      0.534162
        V3      1.008391
        V4      0.313238
        V5      3.326097
        V6     16.635097
        V7      0.541507
        V8      0.700713
        V9      0.123085
        V10     0.597813
        V11     0.918393
        V12     0.201503
        V13     0.493064
        V14    156.100173
        dtype: float64
```

```
In [11]: def printMeanAndSdByGroup(variables, groupvariable):
        data_groupby = variables.groupby(groupvariable)
        print("## Means:")
        display(data_groupby.apply(np.mean))
        print("\n## Standard deviations:")
        display(data_groupby.apply(np.std))
        print("\n## Sample sizes:")
        display(pd.DataFrame(data_groupby.apply(len)))
```

```
In [12]: printMeanAndSdByGroup(X, y)
```

```
## Means:
```

```
V1
```

```
1    98.081473
2    51.077883
3    60.259487
dtype: float64
```

```
## Standard deviations:
```

	V2	V3	V4	V5	V6	V7	V8	V9	V10
V1									
1	0.458192	0.682689	0.225233	2.524651	10.409595	0.336077	0.394111	0.069453	0.408602
2	0.534162	1.008391	0.313238	3.326097	16.635097	0.541507	0.700713	0.123085	0.597813
3	0.524689	1.076514	0.182756	2.234515	10.776433	0.353233	0.290431	0.122840	0.404555

```
## Sample sizes:
```

```
0
```

```
V1
```

```
1  59
2  71
3  48
```

```
In [13]: corr = stats.pearsonr(X.V2, X.V3)
print("p-value:\t", corr[1])
print("cor:\t\t", corr[0])
```


```
p-value:      0.21008198597074274
cor:          0.09439694091041399
```

```
In [14]: corrmat = X.corr()
corrmat
```

```
Out[14]:
```

	V2	V3	V4	V5	V6	V7	V8	V9	
V2	1.000000	0.094397	0.211545	-0.310235	0.270798	0.289101	0.236815	-0.155929	0.136698
V3	0.094397	1.000000	0.164045	0.288500	-0.054575	-0.335167	-0.411007	0.292977	-0.220746
V4	0.211545	0.164045	1.000000	0.443367	0.286587	0.128980	0.115077	0.186230	0.009652
V5	-0.310235	0.288500	0.443367	1.000000	-0.083333	-0.321113	-0.351370	0.361922	-0.197327
V6	0.270798	-0.054575	0.286587	-0.083333	1.000000	0.214401	0.195784	-0.256294	0.236441
V7	0.289101	-0.335167	0.128980	-0.321113	0.214401	1.000000	0.864564	-0.449935	0.612413
V8	0.236815	-0.411007	0.115077	-0.351370	0.195784	0.864564	1.000000	-0.537900	0.652692
V9	-0.155929	0.292977	0.186230	0.361922	-0.256294	-0.449935	-0.537900	1.000000	-0.365845
V10	0.136698	-0.220746	0.009652	-0.197327	0.236441	0.612413	0.652692	-0.365845	1.000000
V11	0.546364	0.248985	0.258887	0.018732	0.199950	-0.055136	-0.172379	0.139057	-0.025136
V12	-0.071747	-0.561296	-0.074667	-0.273955	0.055398	0.433681	0.543479	-0.262640	0.295398
V13	0.072343	-0.368710	0.003911	-0.276769	0.066004	0.699949	0.787194	-0.503270	0.519949
V14	0.643720	-0.192011	0.223626	-0.440597	0.393351	0.498115	0.494193	-0.311385	0.330597

```
In [15]: sns.heatmap(corrmat, vmax=1., square=False).xaxis.tick_top()
```



```
In [17]: def mosthighlycorrelated(mydataframe, numtoreport):
# find the correlations
cormatrix = mydataframe.corr()
# set the correlations on the diagonal or lower triangle to zero,
# so they will not be reported as the highest ones:
cormatrix *= np.tri(*cormatrix.values.shape, k=-1).T
# find the top n correlations
cormatrix = cormatrix.stack()
cormatrix = cormatrix.reindex(cormatrix.abs().sort_values(ascending=False).index)
# assign human-friendly names
cormatrix.columns = ["FirstVariable", "SecondVariable", "Correlation"]
return cormatrix.head(numtoreport)
```

```
In [18]: mosthighlycorrelated(X, 10)
```

```
Out[18]:
```

	FirstVariable	SecondVariable	Correlation
0	V7	V8	0.864564
1	V8	V13	0.787194
2	V7	V13	0.699949
3	V8	V10	0.652692
4	V2	V14	0.643720
5	V7	V10	0.612413
6	V12	V13	0.565468
7	V3	V12	-0.561296
8	V2	V11	0.546364
9	V8	V12	0.543479

```
In [33]: standardisedX = scale(X)
standardisedX = pd.DataFrame(standardisedX, index=X.index, columns=X.columns)
```

```
In [34]: pca = PCA().fit(standardisedX)
```

```
In [35]: def pca_summary(pca, standardised_data, out=True):
names = ["PC"+str(i) for i in range(1, len(pca.explained_variance_ratio_)+1)]
a = list(np.std(pca.transform(standardised_data), axis=0))
b = list(pca.explained_variance_ratio_)
c = [np.sum(pca.explained_variance_ratio_[0:i]) for i in range(1, len(pca.explained_variance_ratio_)+1)]
columns = pd.MultiIndex.from_tuples([("sdev", "Standard deviation"), ("varpro", "variance proportion")])
summary = pd.DataFrame(zip(a, b, c), index=names, columns=columns)
if out:
    print("Importance of components:")
    display(summary)
return summary
```

```
In [36]: summary = pca_summary(pca, standardisedX)
```

Importance of components:

	sdev	varprop	cumprop
	Standard deviation	Proportion of Variance	Cumulative Proportion
PC1	2.169297	0.361988	0.361988
PC2	1.580182	0.192075	0.554063
PC3	1.202527	0.111236	0.665300
PC4	0.958631	0.070690	0.735990
PC5	0.923704	0.065633	0.801623
PC6	0.801035	0.049358	0.850981
PC7	0.742313	0.042387	0.893368
PC8	0.590337	0.026807	0.920175
PC9	0.537476	0.022222	0.942397
PC10	0.500902	0.019300	0.961697
PC11	0.475172	0.017368	0.979066
PC12	0.410817	0.012982	0.992048
PC13	0.321524	0.007952	1.000000

```
In [41]: summary.sdev
```

Out[41]:

	Standard deviation
PC1	2.169297
PC2	1.580182
PC3	1.202527
PC4	0.958631
PC5	0.923704
PC6	0.801035
PC7	0.742313
PC8	0.590337
PC9	0.537476
PC10	0.500902
PC11	0.475172
PC12	0.410817
PC13	0.321524


```
In [38]: lda = LinearDiscriminantAnalysis().fit(X, y)
```

```
In [39]: def pretty_scalings(lda, X, out=False):
    ret = pd.DataFrame(lda.scalings_, index=X.columns, columns=["LD"+str(i+1)
    for i in range(lda.scalings_.shape[1])])
    if out:
        print("Coefficients of linear discriminants:")
        display(ret)
    return ret
pretty_scalings_ = pretty_scalings(lda, X, out=True)
```

Coefficients of linear discriminants:

	LD1	LD2
V2	0.403400	0.871793
V3	-0.165255	0.305380
V4	0.369075	2.345850
V5	-0.154798	-0.146381
V6	0.002163	-0.000463
V7	-0.618052	-0.032213
V8	1.661191	-0.491998
V9	1.495818	-1.630954
V10	-0.134093	-0.307088
V11	-0.355056	0.253231
V12	0.818036	-1.515634
V13	1.157559	0.051184
V14	0.002691	0.002853

```
In [43]: sns.lmplot("LD1", "LD2", lda_values["x"].join(y), hue="V1", fit_reg=False);
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[43], line 1
----> 1 sns.lmplot("LD1", "LD2", lda_values["x"].join(y), hue="V1", fit_reg=False)

NameError: name 'lda_values' is not defined
```

```
In [ ]:
```

